

Korea Advanced Institute of Science and Technology

School of Electrical Engineering

EE817 GPU Programming and Its Application Spring 2018

Student's Name: Dinh Vu

Student's ID: 20184187

Homework 4

The computer, used in my homework 4, contains NVIDIA GeForce GT 1070 based on Pascal GP104 architecture.

```
20184187@eelab5:~/gpu_programming/hw/hw4$ nvidia-smi
Tue May  1 15:33:37 2018
```

NVIDIA-SMI 390.48				Driver Version: 390.48			
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute	M.
0	GeForce GTX 1070	Off	00000000:01:00.0	Off			N/A
0%	44C	P5	12W / 200W	0MiB / 8119MiB	0%	Default	
1	GeForce GTX 1070	Off	00000000:02:00.0	Off			N/A
0%	48C	P5	22W / 200W	0MiB / 8119MiB	2%	Default	

Processes:				GPU Memory
GPU	PID	Type	Process name	Usage
No running processes found				

Figure 1. Graphic card information

1. Problem 7

The source code for problem 7 is prob7.cu file and the results are shown in Figure 1.1 and Table 1.1 below with thread block 32×16 using 4-byte access mode. There is not any load transaction from global memory, so Global Load Transactions Per Request equals to 0.

By padding $IPAD = 2$ elements in each row of the tile in the kernel setRowReadColPad, the even-column elements and the odd-column elements are distributed among even banks and odd banks, respectively. So, both writing by row-major and reading by column-major are conflict-free. It means the number of transactions for a request from shared memory is 1.00.

However, with the same padding, the kernel setColReadRowPad has a 2-way bank conflict in both reading and writing operation. The mapping from words to banks is illustrated at Figure 1.2 and Figure 1.3.

```

20184187@ee1ab5:~/gpu_programming/hw/hw4$ nvprof --metrics gld_transactions_per_request,gst_transactions_per_request,shared_load_transactions_per_request,shared_store_transactions_per_request ./prob7
==17437== Nvprof is profiling process 17437, command: ./prob7
./prob7 at device 0: GeForce GTX 1070 with Bank Model: 4-byte <<grid(1, 1), block(32, 16)>>>
==17437== Some kernel(s) will be replayed on device 0 in order to collect all events/metrics.
==17437== Replaying kernel "setRowReadColPad(int*)" (done)
==17437== Replaying kernel "setColReadRowPad(int*)" (done)
==17437== Profiling application: ./prob7
==17437== Profiling result:
==17437== Metric result:
Invocations
Device "GeForce GTX 1070 (0)"
Kernel: setColReadRowPad(int*)
Metric Name                                Metric Description                                Min           Max           Avg
1 gld_transactions_per_request              Global Load Transactions Per Request                0.000000      0.000000      0.000000
1 gst_transactions_per_request              Global Store Transactions Per Request                4.000000      4.000000      4.000000
1 shared_load_transactions_per_request       Shared Memory Load Transactions Per Request          2.000000      2.000000      2.000000
1 shared_store_transactions_per_request      Shared Memory Store Transactions Per Request          2.000000      2.000000      2.000000
Kernel: setRowReadColPad(int*)
1 gld_transactions_per_request              Global Load Transactions Per Request                0.000000      0.000000      0.000000
1 gst_transactions_per_request              Global Store Transactions Per Request                4.000000      4.000000      4.000000
1 shared_load_transactions_per_request       Shared Memory Load Transactions Per Request          1.000000      1.000000      1.000000
1 shared_store_transactions_per_request      Shared Memory Store Transactions Per Request          1.000000      1.000000      1.000000

```

Figure 1.1. Results of problem 7

Table 1.1. Transaction Metrics

Kernel	setColReadRowPad	setRowReadColPad
Global Load Transactions Per Request	0.00	0.00
Global Store Transactions Per Request	4.00	4.00
Shared Memory Load Transactions Per Request	2.00	1.00
Shared Memory Store Transactions Per Request	2.00	1.00

Bank	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
			32	33	34	35	36	37	38	39	40	41	42	43	44	45	46	47	48	49	50	51	52	53	54	55	56	57	58	59	60	61
	62	63			64	65	66	67	68	69	70	71	72	73	74	75	76	77	78	79	80	81	82	83	84	85	86	87	88	89	90	91
	92	93	94	95			96	97	98	99	100	101	102	103	104	105	106	107	108	109	110	111	112	113	114	115	116	117	118	119	120	121
	122	123	124	125	126	127			128	129	130	131	132	133	134	135	136	137	138	139	140	141	142	143	144	145	146	147	148	149	150	151
	152	153	154	155	156	157	158	159			160	161	162	163	164	165	166	167	168	169	170	171	172	173	174	175	176	177	178	179	180	181
	182	183	184	185	186	187	188	189	190	191			192	193	194	195	196	197	198	199	200	201	202	203	204	205	206	207	208	209	210	211
	212	213	214	215	216	217	218	219	220	221	222	223			224	225	226	227	228	229	230	231	232	233	234	235	236	237	238	239	240	241
	242	243	244	245	246	247	248	249	250	251	252	253	254	255			256	257	258	259	260	261	262	263	264	265	266	267	268	269	270	271
	272	273	274	275	276	277	278	279	280	281	282	283	284	285	286	287			288	289	290	291	292	293	294	295	296	297	298	299	300	301
	302	303	304	305	306	307	308	309	310	311	312	313	314	315	316	317	318	319			320	321	322	323	324	325	326	327	328	329	330	331
	332	333	334	335	336	337	338	339	340	341	342	343	344	345	346	347	348	349	350	351			352	353	354	355	356	357	358	359	360	361
	362	363	364	365	366	367	368	369	370	371	372	373	374	375	376	377	378	379	380	381	382	383			384	385	386	387	388	389	390	391
	392	393	394	395	396	397	398	399	400	401	402	403	404	405	406	407	408	409	410	411	412	413	414	415			416	417	418	419	420	421
	422	423	424	425	426	427	428	429	430	431	432	433	434	435	436	437	438	439	440	441	442	443	444	445	446	447			448	449	450	451
	452	453	454	455	456	457	458	459	460	461	462	463	464	465	466	467	468	469	470	471	472	473	474	475	476	477	478	479			480	481
	482	483	484	485	486	487	488	489	490	491	492	493	494	495	496	497	498	499	500	501	502	503	504	505	506	507	508	509	510	511		

Figure 1.2. setRowReadColPad with bank conflict-free

Bank	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
	0	32	64	96	128	160	192	224	256	288	320	352	384	416	448	480			1	33	65	97	129	161	193	225	257	289	321	353	385	417
	449	481			2	34	66	98	130	162	194	226	258	290	322	354	386	418	450	482			3	35	67	99	131	163	195	227	259	291
	323	355	387	419	451	483			4	36	68	100	132	164	196	228	260	292	324	356	388	420	452	484			5	37	69	101	133	165
	197	229	261	293	325	357	389	421	453	485			6	38	70	102	134	166	198	230	262	294	326	358	390	422	454	486			7	39
	71	103	135	167	199	231	263	295	327	359	391	423	455	487			8	40	72	104	136	168	200	232	264	296	328	360	392	424	456	488
		9	41	73	105	137	169	201	233	265	297	329	361	393	425	457	489			10	42	74	106	138	170	202	234	266	298	330	362	
	394	426	458	490			11	43	75	107	139	171	203	235	267	299	331	363	395	427	459	491			12	44	76	108	140	172	204	236
	268	300	332	364	396	428	460	492			13	45	77	109	141	173	205	237	269	301	333	365	397	429	461	493			14	46	78	110
	142	174	206	238	270	302	334	366	398	430	462	494			15	47	79	111	143	175	207	239	271	303	335	367	399	431	463	495		
	16	48	80	112	144	176	208	240	272	304	336	368	400	432	464	496			17	49	81	113	145	177	209	241	273	305	337	369	401	433
	465	497			18	50	82	114	146	178	210	242	274	306	338	370	402	434	466	498			19	51	83	115	147	179	211	243	275	307
	339	371	403	435	467	499			20	52	84	116	148	180	212	244	276	308	340	372	404	436	468	500			21	53	85	117	149	181
	213	245	277	309	341	373	405	437	469	501			22	54	86	118	150	182	214	246	278	310	342	374	406	438	470	502			23	55
	87	119	151	183	215	247	279	311	343	375	407	439	471	503			24	56	88	120	152	184	216	248	280	312	344	376	408	440	472	504
		25	57	89	121	153	185	217	249	281	313	345	377	409	441	473	505			26	58	90	122	154	186	218	250	282	314	346	378	
	410	442	474	506			27	59	91	123	155	187	219	251	283	315	347	379	411	443	475	507			28	60	92	124	156	188	220	252
	284	316	348	380	412	444	476	508			29	61	93	125	157	189	221	253	285	317	349	381	413	445	477	509			30	62	94	126
	158	190	222	254	286	318	350	382	414	446	478	510			31	63	95	127	159	191	223	255	287	319	351	383	415	447	479	511		

Figure 1.3. setColReadRowPad with bank conflict at yellow elements

2. Problem 8

The source code for problem 8 is prob8.cu file and its results are display in Figure 2.1, Figure 2.2, Figure 2.3 and Figure 2.4 corresponding to block sizes of 64, 128, 512 and 1024. The summary of the results is typed in Table 2.1.

```
20184187@eelab5:~/gpu_programming/hw/hw4$ nvcc -ccbin gcc-4.9 -arch=sm_61 -o prob8 prob8.cu
20184187@eelab5:~/gpu_programming/hw/hw4$ nvprof ./prob8 64
==18586== NVPF is profiling process 18586, command: ./prob8 64
./prob8 starting reduction at device 0: GeForce GTX 1070 with array size 4096 grid 64, block 64
CPU reduce: 517140
reduceGmem <<<grid 64, block 64>>>: 517140 OK!
reduceSmem <<<grid 64, block 64>>>: 517140 OK!
reduceSmemDyn <<<grid 64, block 64>>>: 517140 OK!
reduceGmemUnroll <<<grid 16, block 64>>>: 517140 OK!
reduceSmemUnroll <<<grid 16, block 64>>>: 488250 Failed!
reduceSmemUnrollDyn <<<grid 16, block 64>>>: 517140 OK!
reduceNeighboredGmem <<<grid 64, block 64>>>: 517140 OK!
reduceNeighboredSmem <<<grid 64, block 64>>>: 8545 Failed!
==18586== Profiling application: ./prob8 64
==18586== Profiling result:
Time(%) Time Calls Avg Min Max Name
51.39% 24.927us 8 3.1150us 1.8240us 12.127us [CUDA memcpy HtoD]
11.74% 5.6960us 8 712ns 640ns 960ns [CUDA memcpy DtoH]
6.93% 3.3600us 1 3.3600us 3.3600us 3.3600us reduceNeighboredGmem(int*, int*, unsigned int)
5.41% 2.6240us 1 2.6240us 2.6240us 2.6240us reduceNeighboredSmem(int*, int*, unsigned int)
5.08% 2.4640us 1 2.4640us 2.4640us 2.4640us reduceGmem(int*, int*, unsigned int)
4.88% 2.3680us 1 2.3680us 2.3680us 2.3680us reduceGmemUnroll(int*, int*, unsigned int)
4.15% 2.0150us 1 2.0150us 2.0150us 2.0150us reduceSmemUnrollDyn(int*, int*, unsigned int)
3.56% 1.7280us 1 1.7280us 1.7280us 1.7280us reduceSmem(int*, int*, unsigned int)
3.43% 1.6640us 1 1.6640us 1.6640us 1.6640us reduceSmemUnroll(int*, int*, unsigned int)
3.43% 1.6640us 1 1.6640us 1.6640us 1.6640us reduceSmemDyn(int*, int*, unsigned int)

==18586== API calls:
Time(%) Time Calls Avg Min Max Name
68.32% 105.87ms 2 52.937ms 115.06us 105.76ms cudaMalloc
29.89% 46.320ms 1 46.320ms 46.320ms 46.320ms cudaDeviceReset
0.63% 983.44us 182 5.4030us 311ns 299.76us cuDeviceGetAttribute
0.59% 921.36us 16 57.584us 6.5010us 266.35us cudaMemcpy
0.19% 301.55us 1 301.55us 301.55us 301.55us cudaGetDeviceProperties
0.15% 236.68us 2 118.34us 102.51us 134.18us cudaFree
0.10% 152.76us 2 76.377us 71.863us 80.892us cuDeviceTotalMem
0.05% 84.024us 8 10.503us 6.7080us 21.390us cudaLaunch
0.05% 72.026us 2 36.013us 35.432us 36.594us cuDeviceGetName
0.00% 4.5900us 1 4.5900us 4.5900us 4.5900us cudaSetDevice
0.00% 3.9980us 24 166ns 122ns 299ns cudaSetupArgument
0.00% 2.6610us 6 443ns 340ns 646ns cuDeviceGet
0.00% 2.3120us 8 289ns 295ns 638ns cudaConfigureCall
0.00% 2.1230us 3 707ns 327ns 1.4210us cuDeviceGetCount
0.00% 1.6920us 8 211ns 178ns 351ns cudaGetLastError
```

Figure 2.1. Test block size of 64

```
20184187@eelab5:~/gpu_programming/hw/hw4$ nvprof ./prob8 128
==20894== NVPF is profiling process 20894, command: ./prob8 128
./prob8 starting reduction at device 0: GeForce GTX 1070 with array size 4096 grid 32, block 128
CPU reduce: 517140
reduceGmem <<<grid 32, block 128>>>: 517140 OK!
reduceSmem <<<grid 32, block 128>>>: 517140 OK!
reduceSmemDyn <<<grid 32, block 128>>>: 517140 OK!
reduceGmemUnroll <<<grid 8, block 128>>>: 517140 OK!
reduceSmemUnroll <<<grid 8, block 128>>>: 454778 Failed!
reduceSmemUnrollDyn <<<grid 8, block 128>>>: 517140 OK!
reduceNeighboredGmem <<<grid 32, block 128>>>: 517140 OK!
reduceNeighboredSmem <<<grid 32, block 128>>>: 4404 Failed!
==20894== Profiling application: ./prob8 128
==20894== Profiling result:
Time(%) Time Calls Avg Min Max Name
36.21% 23.071us 8 2.8830us 1.8240us 9.9840us [CUDA memcpy HtoD]
33.95% 21.632us 8 2.7040us 640ns 10.784us [CUDA memcpy DtoH]
5.42% 3.4560us 1 3.4560us 3.4560us 3.4560us reduceNeighboredGmem(int*, int*, unsigned int)
4.42% 2.8160us 1 2.8160us 2.8160us 2.8160us reduceGmem(int*, int*, unsigned int)
4.42% 2.8160us 1 2.8160us 2.8160us 2.8160us reduceGmemUnroll(int*, int*, unsigned int)
4.07% 2.5920us 1 2.5920us 2.5920us 2.5920us reduceNeighboredSmem(int*, int*, unsigned int)
3.06% 1.9520us 1 1.9520us 1.9520us 1.9520us reduceSmem(int*, int*, unsigned int)
2.91% 1.8560us 1 1.8560us 1.8560us 1.8560us reduceSmemUnrollDyn(int*, int*, unsigned int)
2.86% 1.8240us 1 1.8240us 1.8240us 1.8240us reduceSmemDyn(int*, int*, unsigned int)
2.66% 1.6960us 1 1.6960us 1.6960us 1.6960us reduceSmemUnroll(int*, int*, unsigned int)

==20894== API calls:
Time(%) Time Calls Avg Min Max Name
70.61% 118.06ms 2 59.031ms 111.96us 117.95ms cudaMalloc
27.80% 46.482ms 1 46.482ms 46.482ms 46.482ms cudaDeviceReset
0.57% 947.97us 182 5.2080us 309ns 281.01us cuDeviceGetAttribute
0.46% 764.68us 16 47.792us 6.5350us 260.10us cudaMemcpy
0.18% 294.35us 1 294.35us 294.35us 294.35us cudaGetDeviceProperties
0.14% 235.49us 2 117.74us 113.33us 122.16us cudaFree
0.10% 169.16us 2 84.581us 83.784us 85.378us cuDeviceTotalMem
0.10% 164.91us 8 20.613us 6.9460us 35.983us cudaLaunch
0.04% 72.720us 2 36.360us 34.821us 37.899us cuDeviceGetName
0.00% 4.9590us 8 619ns 198ns 3.3400us cudaConfigureCall
0.00% 4.5140us 1 4.5140us 4.5140us 4.5140us cudaSetDevice
0.00% 3.7520us 24 156ns 116ns 312ns cudaSetupArgument
0.00% 2.5860us 3 862ns 318ns 1.3350us cuDeviceGetCount
0.00% 2.4550us 6 409ns 325ns 539ns cuDeviceGet
0.00% 1.7660us 8 220ns 178ns 266ns cudaGetLastError
```

Figure 2.2. Test block size of 128

```

20184187@eelab5:~/gpu_programming/hw/hw4$ nvprof ./prob8 512
==20906== NVPROF is profiling process 20906, command: ./prob8 512
./prob8 starting reduction at device 0: GeForce GTX 1070 with array size 4096 grid 8, block 512
CPU reduce: 517140
reduceGmem <<<grid 8, block 512>>>: 517140 OK!
reduceSmem <<<grid 8, block 512>>>: 517140 OK!
reduceSmemDyn <<<grid 8, block 512>>>: 517140 OK!
reduceGmemUnroll <<<grid 2, block 512>>>: 517140 OK!
reduceSmemUnroll <<<grid 2, block 512>>>: 264675 Failed!
reduceSmemUnrollDyn <<<grid 2, block 512>>>: 517140 OK!
reduceNeighboredGmem <<<grid 8, block 512>>>: 517140 OK!
reduceNeighboredSmem <<<grid 8, block 512>>>: 1106 Failed!
==20906== Profiling application: ./prob8 512
==20906== Profiling result:
Time(%) Time Calls Avg Min Max Name
36.35% 14.527us 8 1.8150us 1.7920us 1.8240us [CUDA memcpy HtoD]
13.53% 5.4080us 8 676ns 640ns 768ns [CUDA memcpy DtoH]
10.89% 4.3520us 1 4.3520us 4.3520us 4.3520us reduceNeighboredGmem(int*, int*, unsigned int)
7.53% 3.0080us 1 3.0080us 3.0080us 3.0080us reduceGmemUnroll(int*, int*, unsigned int)
7.13% 2.8480us 1 2.8480us 2.8480us 2.8480us reduceNeighboredSmem(int*, int*, unsigned int)
7.13% 2.8480us 1 2.8480us 2.8480us 2.8480us reduceGmem(int*, int*, unsigned int)
4.64% 1.8560us 1 1.8560us 1.8560us 1.8560us reduceSmemUnroll(int*, int*, unsigned int)
4.40% 1.7600us 1 1.7600us 1.7600us 1.7600us reduceSmemUnrollDyn(int*, int*, unsigned int)
4.24% 1.6960us 1 1.6960us 1.6960us 1.6960us reduceSmemDyn(int*, int*, unsigned int)
4.16% 1.6640us 1 1.6640us 1.6640us 1.6640us reduceSmem(int*, int*, unsigned int)

==20906== API calls:
Time(%) Time Calls Avg Min Max Name
66.71% 100.59ms 2 50.296ms 186.65us 100.41ms cudaMalloc
31.50% 47.508ms 1 47.508ms 47.508ms 47.508ms cudaDeviceReset
0.63% 954.90us 182 5.2460us 315ns 281.22us cuDeviceGetAttribute
0.57% 855.32us 16 53.457us 6.5990us 209.79us cudaMemcpy
0.20% 304.25us 1 304.25us 304.25us 304.25us cudaGetDeviceProperties
0.17% 255.81us 2 127.91us 112.19us 143.62us cudaFree
0.10% 154.95us 2 77.473us 72.245us 82.701us cuDeviceTotalMem
0.06% 84.830us 8 10.603us 6.5810us 21.808us cudaLaunch
0.05% 70.389us 2 35.194us 35.152us 35.237us cuDeviceGetName
0.00% 4.9710us 1 4.9710us 4.9710us 4.9710us cudaSetDevice
0.00% 3.6830us 24 153ns 119ns 329ns cudaSetupArgument
0.00% 2.6920us 6 448ns 335ns 691ns cuDeviceGet
0.00% 2.4660us 8 308ns 216ns 783ns cudaConfigureCall
0.00% 2.0960us 3 698ns 312ns 1.2850us cuDeviceGetCount
0.00% 1.7590us 8 219ns 174ns 360ns cudaGetLastError

```

Figure 2.3. Test block size of 512

```

20184187@eelab5:~/gpu_programming/hw/hw4$ nvprof ./prob8 1024
==20918== NVPROF is profiling process 20918, command: ./prob8 1024
./prob8 starting reduction at device 0: GeForce GTX 1070 with array size 4096 grid 4, block 1024
CPU reduce: 517140
reduceGmem <<<grid 4, block 1024>>>: 517140 OK!
reduceSmem <<<grid 4, block 1024>>>: 517140 OK!
reduceSmemDyn <<<grid 4, block 1024>>>: 517140 OK!
reduceGmemUnroll <<<grid 1, block 1024>>>: 517140 OK!
reduceSmemUnroll <<<grid 1, block 1024>>>: 522 Failed!
reduceSmemUnrollDyn <<<grid 1, block 1024>>>: 517140 OK!
reduceNeighboredGmem <<<grid 4, block 1024>>>: 517140 OK!
reduceNeighboredSmem <<<grid 4, block 1024>>>: 736 Failed!
==20918== Profiling application: ./prob8 1024
==20918== Profiling result:
Time(%) Time Calls Avg Min Max Name
31.13% 17.824us 8 2.2280us 640ns 11.520us [CUDA memcpy DtoH]
29.68% 16.992us 8 2.1240us 1.7920us 4.3840us [CUDA memcpy HtoD]
9.22% 5.2800us 1 5.2800us 5.2800us 5.2800us reduceNeighboredGmem(int*, int*, unsigned int)
6.93% 3.9680us 1 3.9680us 3.9680us 3.9680us reduceNeighboredSmem(int*, int*, unsigned int)
5.20% 2.9760us 1 2.9760us 2.9760us 2.9760us reduceGmemUnroll(int*, int*, unsigned int)
4.97% 2.8480us 1 2.8480us 2.8480us 2.8480us reduceGmem(int*, int*, unsigned int)
3.52% 2.0160us 1 2.0160us 2.0160us 2.0160us reduceSmemUnrollDyn(int*, int*, unsigned int)
3.19% 1.8240us 1 1.8240us 1.8240us 1.8240us reduceSmem(int*, int*, unsigned int)
3.19% 1.8240us 1 1.8240us 1.8240us 1.8240us reduceSmemUnroll(int*, int*, unsigned int)
2.96% 1.6960us 1 1.6960us 1.6960us 1.6960us reduceSmemDyn(int*, int*, unsigned int)

==20918== API calls:
Time(%) Time Calls Avg Min Max Name
66.81% 100.29ms 2 50.144ms 111.21us 100.18ms cudaMalloc
31.56% 47.379ms 1 47.379ms 47.379ms 47.379ms cudaDeviceReset
0.66% 988.15us 182 5.4290us 307ns 279.15us cuDeviceGetAttribute
0.39% 592.90us 16 37.056us 6.4820us 128.55us cudaMemcpy
0.20% 295.50us 1 295.50us 295.50us 295.50us cudaGetDeviceProperties
0.14% 212.95us 2 106.47us 100.16us 112.79us cudaFree
0.11% 161.72us 2 80.861us 78.954us 82.769us cuDeviceTotalMem
0.07% 103.83us 8 12.978us 6.4370us 28.800us cudaLaunch
0.05% 68.886us 2 34.443us 33.815us 35.071us cuDeviceGetName
0.00% 4.7010us 1 4.7010us 4.7010us 4.7010us cudaSetDevice
0.00% 3.9630us 24 165ns 122ns 301ns cudaSetupArgument
0.00% 2.9510us 8 368ns 238ns 1.1550us cudaConfigureCall
0.00% 2.3850us 6 397ns 313ns 553ns cuDeviceGet
0.00% 2.0020us 3 667ns 344ns 1.2820us cuDeviceGetCount
0.00% 1.7380us 8 217ns 171ns 380ns cudaGetLastError

```

Figure 2.4. Test block size of 1024

Table 2.1. Performance on various block sizes

Kernel	Execution time in μs			
	64 threads	128 threads	512 threads	1024 threads
reduceGmem	2.4640	2.8160	2.8480	2.8480
reduceSmem	1.7280	1.9520	1.6640	1.8240
reduceSmemDyn	1.6640	1.8240	1.6960	1.6960
reduceGmemUnroll	2.3680	2.8160	3.0080	2.9760
reduceSmemUnroll	1.6640	1.6960	1.8560	1.8240
reduceSmemUnrollDyn	2.0150	1.8560	1.7600	2.0160
reduceNeighboredGmem	3.3600	3.4560	4.3520	5.2800
reduceNeighboredSmem	2.6240	2.5920	2.8480	3.9680

In general, with Pascal architecture, the best execution configuration is 64 threads per block, though reduceSmem, reduceSmemDyn and reduceSmemUnrollDyn achieve the best performance with 512 threads per block. The execution time with block size of 1024 is always highest.

3. Problem 11

For problem 11, the source code is prob11.cu file and the results, depicted in Figure 3.1 below. Based on the kernel test_shfl_wrap, the kernel test_shfl_wrap_plus increases the current thread's value by the value of the thread, which is two indexes greater. Because the width of the shuffle operation is still 16, so the first two initial threads warp around to the bottom two threads. The operation of the kernel test_shfl_wrap_plus, presented in Figure 3.2.


```

20184187@eelab5:~/gpu_programming/hw/hw4$ nvcc -ccbin gcc-4.9 -arch=sm_61 -o prob11 prob11.cu
20184187@eelab5:~/gpu_programming/hw/hw4$ nvprof ./prob11
==21076== NVPROF is profiling process 21076, command: ./prob11
> ./prob11 Starting at Device 0: GeForce GTX 1070
Initial data      : 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15
test_shfl_wrap    : 14 15 0  1  2  3  4  5  6  7  8  9 10 11 12 13
test_shfl_wrap_plus : 2  4  6  8 10 12 14 16 18 20 22 24 26 28 14 16
==21076== Profiling application: ./prob11
==21076== Profiling result:
Time(%)    Time      Calls      Avg      Min      Max  Name
69.33%    8.8960us      2  4.4480us  640ns  8.2560us  [CUDA memcpy DtoH]
13.47%    1.7280us      1  1.7280us  1.7280us  1.7280us  test_shfl_wrap(int*, int*, int)
10.97%    1.4080us      1  1.4080us  1.4080us  1.4080us  test_shfl_wrap_plus(int*, int*, int)
 6.23%      800ns      1    800ns    800ns    800ns    [CUDA memcpy HtoD]

==21076== API calls:
Time(%)    Time      Calls      Avg      Min      Max  Name
67.99%   100.42ms      2  50.210ms  8.4110us  100.41ms  cudaMalloc
30.79%    45.485ms      1  45.485ms  45.485ms  45.485ms  cudaDeviceReset
 0.64%    938.65us    182  5.1570us   317ns  275.19us  cuDeviceGetAttribute
 0.23%    332.58us      1  332.58us  332.58us  332.58us  cudaGetDeviceProperties
 0.10%    153.80us      2   76.897us  22.591us  131.20us  cudaFree
 0.10%    147.09us      2   73.546us  68.930us  78.163us  cuDeviceTotalMem
 0.07%    105.77us      3   35.257us  10.808us  52.567us  cudaMemcpy
 0.05%    69.503us      2   34.751us  34.451us  35.052us  cuDeviceGetName
 0.02%    24.141us      2   12.070us  8.5380us  15.603us  cudaLaunch
 0.01%    22.084us      1   22.084us  22.084us  22.084us  cudaSetDevice
 0.00%    2.5510us      6    425ns   323ns    567ns  cuDeviceGet
 0.00%    2.1360us      3    712ns   332ns   1.3950us  cuDeviceGetCount
 0.00%    1.1800us      6    196ns   117ns    291ns  cudaSetupArgument
 0.00%    1.0250us      2    512ns   379ns    646ns  cudaConfigureCall
 0.00%      919ns      2    459ns   223ns    696ns  cudaGetLastError

```

Figure 3.1. The results of problem 11

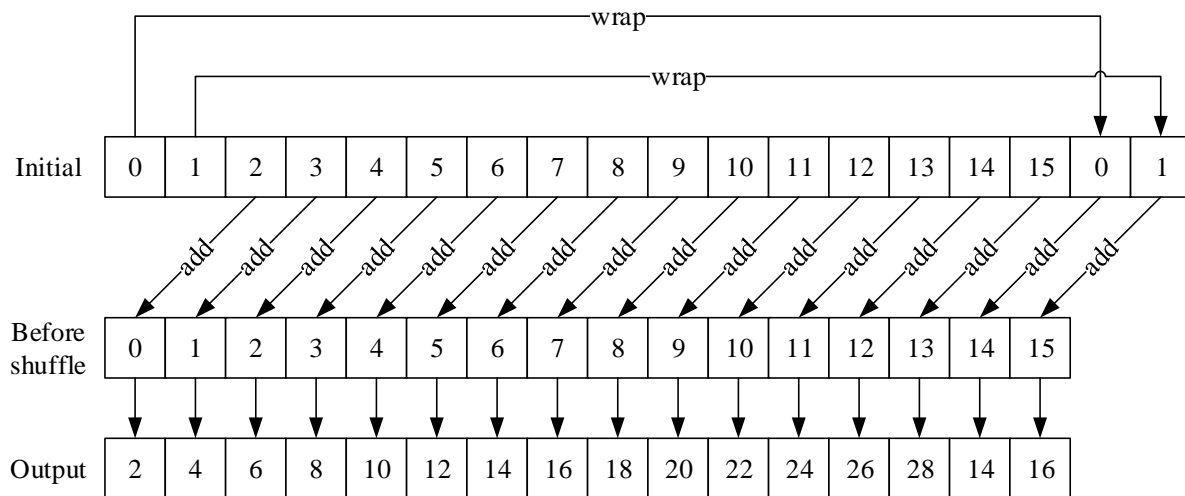


Figure 3.2. The operation of the kernel `test_shfl_wrap_plus`

4. Problem 12

The `prob12.cu` file is the source code of problem 12 and its results are shown in Figure 4.1. Based on the kernel `test_shfl_xor`, the kernel named `test_shfl_xor_plus` simply performs `__shfl_xor` function by passing one as the mask and increases the current thread's value with the received value. Every even thread adds the value of the odd thread above it and every odd thread receives the value of the even thread below it.

```

20184187@eelab5:~/gpu_programming/hw/hw4$ nvcc -ccbin gcc-4.9 -arch=sm_61 -o prob12 prob12.cu
20184187@eelab5:~/gpu_programming/hw/hw4$ nvprof ./prob12
==21678== NVPROF is profiling process 21678, command: ./prob12
> ./prob12 Starting at Device 0: GeForce GTX 1070
Initial Data      : 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15
test_shfl_xor     :  1  0  3  2  5  4  7  6  9  8 11 10 13 12 15 14
test_shfl_xor_plus:  1  1  5  5  9  9 13 13 17 17 21 21 25 25 29 29
==21678== Profiling application: ./prob12
==21678== Profiling result:
Time(%)    Time      Calls      Avg      Min      Max  Name
33.52%    1.9200us      1  1.9200us  1.9200us  1.9200us  test_shfl_xor(int*, int*, int)
28.49%    1.6320us      1  1.6320us  1.6320us  1.6320us  test_shfl_xor_plus(int*, int*, int)
25.14%    1.4400us      2    720ns    704ns    736ns  [CUDA memcpy DtoH]
12.85%     736ns        1    736ns    736ns    736ns  [CUDA memcpy HtoD]

==21678== API calls:
Time(%)    Time      Calls      Avg      Min      Max  Name
75.58%    146.20ms      2  73.099ms  6.6360us  146.19ms  cudaMalloc
23.20%    44.868ms      1  44.868ms  44.868ms  44.868ms  cudaDeviceReset
0.68%     1.3071ms    182  7.1810us   324ns   304.52us  cuDeviceGetAttribute
0.32%     610.88us      1  610.88us  610.88us  610.88us  cudaGetDeviceProperties
0.08%     163.66us      2  81.828us  80.736us  82.921us  cuDeviceTotalMem
0.06%     115.14us      2  57.568us  10.485us  104.65us  cudaFree
0.04%     81.525us      2  40.762us  37.393us  44.132us  cuDeviceGetName
0.02%     40.626us      3  13.542us  11.612us  16.743us  cudaMemcpy
0.02%     35.730us      2  17.865us  14.856us  20.874us  cudaLaunch
0.00%     7.3180us      1  7.3180us  7.3180us  7.3180us  cudaSetDevice
0.00%     2.8270us      6    471ns    349ns    697ns  cuDeviceGet
0.00%     2.0790us      3    693ns    327ns   1.2660us  cuDeviceGetCount
0.00%     1.3120us      6    218ns    135ns    341ns  cudaSetupArgument
0.00%      985ns      2    492ns    337ns    648ns  cudaConfigureCall
0.00%      760ns      2    380ns    355ns    405ns  cudaGetLastError

```

Figure 4.1. The results of problem 12

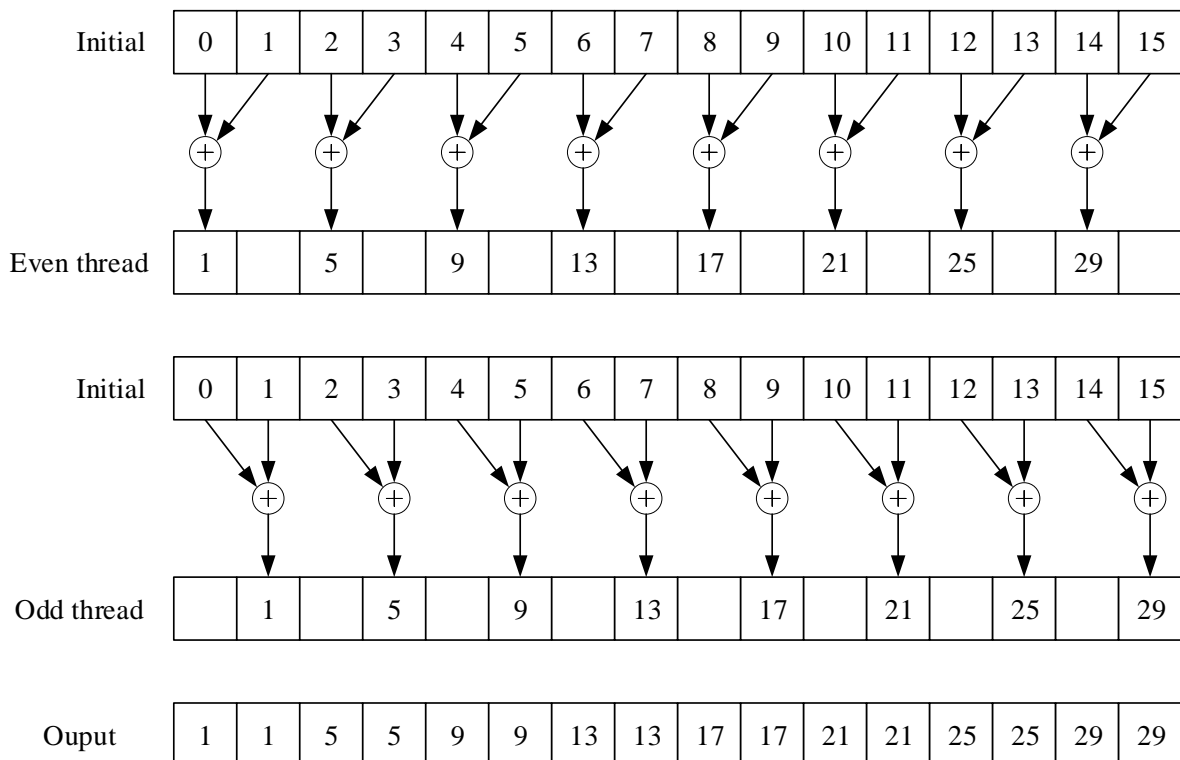


Figure 4.2. The operation of the kernel `test_shfl_xor_plus`