

Student's Name: Dinh Vu

Student's ID: 20184187

Homework 5

The computer, used in my homework 4, contains NVIDIA GeForce GT 1070 based on Pascal GP104 architecture.

```
20184187@eelab5:~/gpu_programming/hw/hw5$ nvidia-smi
Thu May 10 15:56:20 2018

+-----+
| NVIDIA-SMI 390.48                  Driver Version: 390.48          |
+-----+-----+
| GPU   Name           Persistence-M| Bus-Id        Disp.A | Volatile Uncorr. ECC |
| Fan  Temp  Perf    Pwr:Usage/Cap|      Memory-Usage | GPU-Util  Compute M. |
+-----+-----+
|  0   GeForce GTX 1070      Off   | 00000000:01:00.0 Off |           N/A       |
| 43%   69C    P2      96W / 200W | 331MiB / 8119MiB |      71%    Default  |
+-----+-----+
|  1   GeForce GTX 1070      Off   | 00000000:02:00.0 Off |           N/A       |
| 1%    45C    P8      11W / 200W | 10MiB / 8119MiB |       0%    Default  |
+-----+-----+

+-----+
| Processes:                         GPU Memory |
|  GPU       PID    Type    Process name      Usage      |
+-----+-----+
|    0      28676     C   .../local/MATLAB/R2017a/bin/glnxa64/MATLAB  319MiB |
+-----+-----+
```

Figure 1.1. Graphic card information

1. Without Stream

The source code for matrix multiplication using only global memory without stream is `matrixMulGmem.cu` file.

Each thread computes value of each element in matrix C. Each row of matrix A and the corresponding column of matrix B is read from global memory. Then each element in matrix C with thread index is calculated parallel. This programming strategy is presented in Figure 1.2.

Figure 1.3 shows the procedure of the program while Figure 1.4 displays the summary of the execution time.




Figure 1.2. Matrix multiplication only using global memory

```

20184187@eeelab5:~/gpu_programming/hw/hw5$ nvprof --print-gpu-trace ./gmem
==6566== NVPROF is profiling process 6566, command: ./gmem
Matrix Multiplication OK!
==6566== Profiling application: ./gmem
==6566== Profiling result:
  Start Duration   Grid Size   Block Size   Regs*   SSMem*   DSMem*   Size   Throughput   Device   Context   Stream   Name
  ----  -
252.62ms 1.7475ms      (50 100 1)   (32 16 1)    25      0B      0B      9.7656MB 5.4574GB/s   GeForce GTX 107 1 7 [CUDA memcpy HtoD]
254.46ms 1.7284ms      (50 100 1)   (32 16 1)    25      0B      0B      9.7656MB 5.5178GB/s   GeForce GTX 107 1 7 [CUDA memcpy HtoD]
256.74ms 29.745ms      (50 100 1)   (32 16 1)    25      0B      0B      9.7656MB 5.7666GB/s   GeForce GTX 107 1 7 matrixMulGmem(int*, int*, int*, int) [209]
286.49ms 1.6538ms      (50 100 1)   (32 16 1)    25      0B      0B      9.7656MB 5.7666GB/s   GeForce GTX 107 1 7 [CUDA memcpy DtoH]

```

Figure 1.3. The procedure of the program using only global memory

```

20184187@eelab5:~/gpu_programming/hw/hw5$ nvcc -ccbin gcc-4.9 -arch=sm_61 -o gmem matrixMulGmem.cu
20184187@eelab5:~/gpu_programming/hw/hw5$ nvprof ./gmem
==6542== NVPROF is profiling process 6542, command: ./gmem
Matrix Multiplication OK!
==6542== Profiling application: ./gmem
==6542== Profiling result:
Time(%)   Time      Calls      Avg      Min      Max      Name
85.45%    30.096ms      1    30.096ms  30.096ms  30.096ms  matrixMulGmem(int*, int*, int*, int)
9.74%     3.4295ms      2     1.7147ms  1.7021ms  1.7274ms  [CUDA memcpy HtoD]
4.81%     1.6940ms      1     1.6940ms  1.6940ms  1.6940ms  [CUDA memcpy DtoH]

==6542== API calls:
Time(%)   Time      Calls      Avg      Min      Max      Name
73.07%    105.40ms      2    52.700ms  1.4130us  105.40ms  cudaEventCreate
25.20%    36.347ms      3    12.116ms  1.8342ms  32.118ms  cudaMemcpy
0.65%     934.01us    182    5.1310us   320ns    274.35us  cuDeviceGetAttribute
0.58%     834.19us      3    278.06us  204.00us  318.11us  cudaFree
0.32%     454.61us      3    151.54us  116.08us  215.93us  cudaMalloc
0.11%     153.26us      2     76.628us  73.182us  80.075us  cuDeviceTotalMem
0.05%     78.936us      2     39.468us  38.744us  40.192us  cuDeviceGetName
0.02%     21.709us      1     21.709us  21.709us  21.709us  cudaLaunch
0.01%     11.905us      2     5.9520us  5.4080us  6.4970us  cudaEventRecord
0.00%      3.4580us      4         864ns    156ns    2.7960us  cudaSetupArgument
0.00%      2.8170us      1     2.8170us  2.8170us  2.8170us  cudaEventSynchronize
0.00%      2.7410us      2     1.3700us   806ns    1.9350us  cudaEventDestroy
0.00%      2.6460us      6         441ns    322ns     702ns  cuDeviceGet
0.00%      2.3880us      3         796ns    327ns    1.6850us  cuDeviceGetCount
0.00%      2.0320us      1     2.0320us  2.0320us  2.0320us  cudaEventElapsedTime
0.00%      1.2280us      1     1.2280us  1.2280us  1.2280us  cudaConfigureCall

```

Figure 1.4. The execution time of the matrix multiplication using only global memory

2. Depth-first

The `matrixMulDepth.cu` is source code for matrix multiplication overlapping kernel execution with depth-first data transfer. Because I remote the computer in Haedong Lounge, **nvvp** cannot display the profile of my program. Therefore, instead using **nvvp**, the option `--print-gpu-trace` of **nvprof** command is used to display the sequence of streams.

Matrix B is still load entirely from global memory. Because the required number of streams is 8, matrix A is divided to 8 slides in order to split into each stream. Each stream will read each slide $A[i]$ from global memory then calculate the output slide $C[i]$, where $i = 0, 1, 2, \dots, 8$. Hence, each slide of the output matrix C also is computed following each stream as Figure 2.1.

All streams using the same kernel `matrixMulDepth`. The programming strategy of kernel `matrixMulDepth` is presented in Figure 2.2.

The order operation of the streams and the summary of execution time of each function are shown in Figure 2.3 and Figure 2.4, respectively.

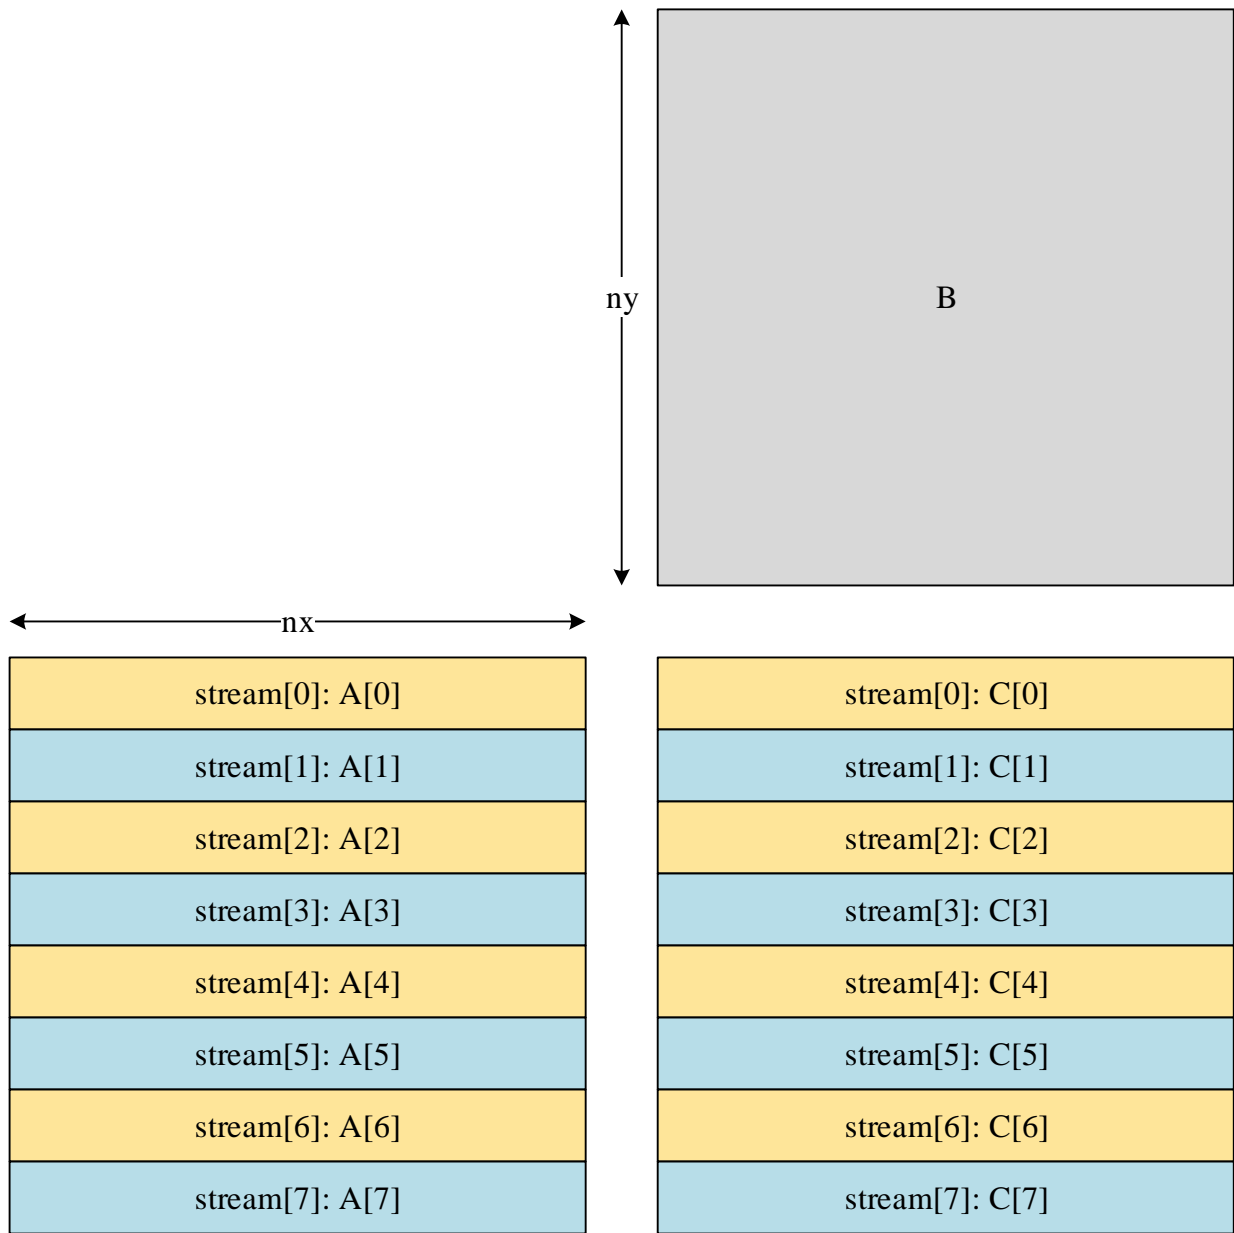


Figure 2.1. The programming strategy for matrix multiplication using streams




Figure 2.2. The operation of kernel `matrixMulDepth`

```
20184187@ee1ab5:~/gpu_programming/hw/hws$ nvprof --print-gpu-trace ./depth
==6669== NVPROF is profiling process 6669, command: ./depth
Matrix Multiplication OK!
==6669== Profiling application: ./depth
==6669== Profiling result:
```

| Start | Duration | Grid Size | Block Size | Regs* | SSMem* | DSMem* | Size | Throughput | Device | Context | Stream | Name |
|----------|----------|-----------|------------|-------|--------|--------|----------|------------|-----------------|---------|--------|---|
| 246.08ms | 213.60us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.5811GB/s | GeForce GTX 107 | 1 | 14 | [CUDA memcpy HtoD] |
| 248.39ms | 452.29ms | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.8778GB/s | GeForce GTX 107 | 1 | 14 | matrixMulDepth(int*, int*, int*, int) [218] |
| 700.69ms | 202.81us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.5231GB/s | GeForce GTX 107 | 1 | 14 | [CUDA memcpy DtoH] |
| 700.92ms | 215.84us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.5231GB/s | GeForce GTX 107 | 1 | 15 | [CUDA memcpy HtoD] |
| 701.17ms | 454.00ms | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.8510GB/s | GeForce GTX 107 | 1 | 15 | matrixMulDepth(int*, int*, int*, int) [227] |
| 1.15517s | 203.74us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.8510GB/s | GeForce GTX 107 | 1 | 15 | [CUDA memcpy DtoH] |
| 1.15540s | 213.28us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.5894GB/s | GeForce GTX 107 | 1 | 16 | [CUDA memcpy HtoD] |
| 1.15593s | 449.94ms | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.8658GB/s | GeForce GTX 107 | 1 | 16 | matrixMulDepth(int*, int*, int*, int) [236] |
| 1.60580s | 203.23us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.6012GB/s | GeForce GTX 107 | 1 | 16 | [CUDA memcpy DtoH] |
| 1.60611s | 212.83us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.6012GB/s | GeForce GTX 107 | 1 | 17 | [CUDA memcpy HtoD] |
| 1.60709s | 454.61ms | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.8843GB/s | GeForce GTX 107 | 1 | 17 | matrixMulDepth(int*, int*, int*, int) [245] |
| 2.06170s | 202.59us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.5544GB/s | GeForce GTX 107 | 1 | 17 | [CUDA memcpy DtoH] |
| 2.06193s | 214.62us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.5544GB/s | GeForce GTX 107 | 1 | 18 | [CUDA memcpy HtoD] |
| 2.06218s | 451.66ms | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.8538GB/s | GeForce GTX 107 | 1 | 18 | matrixMulDepth(int*, int*, int*, int) [254] |
| 2.51305s | 203.65us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.5944GB/s | GeForce GTX 107 | 1 | 18 | [CUDA memcpy DtoH] |
| 2.51408s | 213.09us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.5944GB/s | GeForce GTX 107 | 1 | 19 | [CUDA memcpy HtoD] |
| 2.51433s | 455.14ms | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.8400GB/s | GeForce GTX 107 | 1 | 19 | matrixMulDepth(int*, int*, int*, int) [263] |
| 2.96948s | 204.13us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.5223GB/s | GeForce GTX 107 | 1 | 19 | [CUDA memcpy DtoH] |
| 2.96971s | 215.87us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.8658GB/s | GeForce GTX 107 | 1 | 20 | [CUDA memcpy HtoD] |
| 2.96997s | 452.55ms | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.8658GB/s | GeForce GTX 107 | 1 | 20 | matrixMulDepth(int*, int*, int*, int) [272] |
| 3.42253s | 203.23us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.1363GB/s | GeForce GTX 107 | 1 | 20 | [CUDA memcpy DtoH] |
| 3.42275s | 232.09us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.1363GB/s | GeForce GTX 107 | 1 | 21 | [CUDA memcpy HtoD] |
| 3.42407s | 454.60ms | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.9274GB/s | GeForce GTX 107 | 1 | 21 | matrixMulDepth(int*, int*, int*, int) [281] |
| 3.87868s | 201.12us | (50 13 1) | (32 16 1) | 25 | 0B | 0B | 1.2207MB | 5.9274GB/s | GeForce GTX 107 | 1 | 21 | [CUDA memcpy DtoH] |

Figure 2.3. The procedure of streams in matrix multiplication with depth-approach

```

20184187@eelab5:~/gpu_programming/hw/hw5$ nvprof ./depth
==6637== NVPROF is profiling process 6637, command: ./depth
Matrix Multiplication OK!
==6637== Profiling application: ./depth
==6637== Profiling result:
Time(%)      Time      Calls      Avg      Min      Max      Name
99.91%    3.62599s        8    453.25ms  450.24ms  455.69ms  matrixMulDepth(int*, int*, int*, int)
 0.05%    1.7240ms        8    215.50us  213.18us  223.48us  [CUDA memcpy HtoD]
 0.04%    1.6189ms        8    202.36us  198.62us  204.99us  [CUDA memcpy DtoH]

==6637== API calls:
Time(%)      Time      Calls      Avg      Min      Max      Name
97.03%    3.63279s        8    454.10ms  450.96ms  456.14ms  cudaStreamSynchronize
 2.70%    101.23ms        8    12.654ms  8.8800us  101.16ms  cudaStreamCreate
 0.21%    7.8678ms        4    1.9669ms  1.9424ms  1.9919ms  cudaHostAlloc
 0.03%    946.84us       182    5.2020us    307ns    280.33us  cuDeviceGetAttribute
 0.01%    400.43us        3    133.48us  111.88us  170.06us  cudaMalloc
 0.00%    181.09us       16    11.318us   6.5330us  17.027us  cudaMemcpyAsync
 0.00%    179.16us        8    22.395us   15.006us  39.706us  cudaLaunch
 0.00%    160.66us        2    80.328us   78.369us  82.288us  cuDeviceTotalMem
 0.00%    69.221us        2    34.610us   34.482us  34.739us  cuDeviceGetName
 0.00%    9.9550us       32     311ns    117ns    3.2800us  cudaSetupArgument
 0.00%    6.6450us        8     830ns    613ns    1.3690us  cudaConfigureCall
 0.00%    2.5440us        1    2.5440us  2.5440us  2.5440us  cudaHostGetDevicePointer
 0.00%    2.4660us        6     411ns    322ns    667ns    cuDeviceGet
 0.00%    2.1170us        3     705ns    316ns    1.3900us  cuDeviceGetCount

```

Figure 2.4. Summary of the activities on GPU of matrix multiplication with depth-approach

3. Breadth-first

The source code for matrix multiplication using breadth-first data transfer is matrixMulBreadth.cu. The programming strategy is similar to the depth-first. The only difference is that each activity copy data from host to device, kernel and copy result from device to host have each for-loop.

```

20184187@eelab5:~/gpu_programming/hw/hw5$ nvprof --print-gpu-trace ./breadth
==6914== NVPROF is profiling process 6914, command: ./breadth
Matrix Multiplication OK!
==6914== Profiling application: ./breadth
==6914== Profiling result:
Start Duration      Grid Size      Block Size      Regs*      SSMem*      DSMem*      Size      Throughput      Device      Context      Stream      Name
246.99ms 213.37us - - - - - - 1.2207MB 5.5809GB/s GeForce GTX 107 1 14 [CUDA memcpy HtoD]
247.21ms 212.54us - - - - - - 1.2207MB 5.6808GB/s GeForce GTX 107 1 18 [CUDA memcpy HtoD]
247.43ms 211.90us - - - - - - 1.2207MB 5.6257GB/s GeForce GTX 107 1 16 [CUDA memcpy HtoD]
247.64ms 211.55us - - - - - - 1.2207MB 5.6351GB/s GeForce GTX 107 1 20 [CUDA memcpy HtoD]
247.86ms 212.73us - - - - - - 1.2207MB 5.6037GB/s GeForce GTX 107 1 17 [CUDA memcpy HtoD]
248.08ms 211.93us - - - - - - 1.2207MB 5.6249GB/s GeForce GTX 107 1 21 [CUDA memcpy HtoD]
248.29ms 212.76us - - - - - - 1.2207MB 5.6829GB/s GeForce GTX 107 1 15 [CUDA memcpy HtoD]
248.51ms 213.76us - - - - - - 1.2207MB 3.7994GB/s GeForce GTX 107 1 19 [CUDA memcpy HtoD]
248.68ms 453.78ms (50 13 1) (32 16 1) 25 0B 0B - - GeForce GTX 107 1 14 matrixMulBreadth(int*, int*, int*, int) [225]
702.54ms 456.02ms (50 13 1) (32 16 1) 25 0B 0B - - GeForce GTX 107 1 15 matrixMulBreadth(int*, int*, int*, int) [232]
1.15905s 453.57ms (50 13 1) (32 16 1) 25 0B 0B - - GeForce GTX 107 1 16 matrixMulBreadth(int*, int*, int*, int) [239]
1.61298s 452.92ms (50 13 1) (32 16 1) 25 0B 0B - - GeForce GTX 107 1 17 matrixMulBreadth(int*, int*, int*, int) [246]
2.06589s 453.14ms (50 13 1) (32 16 1) 25 0B 0B - - GeForce GTX 107 1 18 matrixMulBreadth(int*, int*, int*, int) [253]
2.51910s 456.50ms (50 13 1) (32 16 1) 25 0B 0B - - GeForce GTX 107 1 19 matrixMulBreadth(int*, int*, int*, int) [260]
2.97567s 453.60ms (50 13 1) (32 16 1) 25 0B 0B - - GeForce GTX 107 1 20 matrixMulBreadth(int*, int*, int*, int) [267]
3.43083s 452.09ms (50 13 1) (32 16 1) 25 0B 0B - - GeForce GTX 107 1 21 matrixMulBreadth(int*, int*, int*, int) [274]

```

Figure 3.1. The procedure of streams in matrix multiplication with bread-approach

```

20184187@eelab5:~/gpu_programming/hw/hw5$ nvprof ./breadth
==6894== NVPROF is profiling process 6894, command: ./breadth
Matrix Multiplication OK!
==6894== Profiling application: ./breadth
==6894== Profiling result:
Time(%)      Time      Calls      Avg      Min      Max      Name
99.92%    3.61648s        8    452.06ms    449.44ms    454.90ms    matrixMulBreadth(int*, int*, int*, int)
 0.08%    2.9011ms        8    362.64us    211.10us    660.50us    [CUDA memcpy HtoD]

==6894== API calls:
Time(%)      Time      Calls      Avg      Min      Max      Name
96.97%    3.62068s        8    452.59ms    449.72ms    454.94ms    cudaStreamSynchronize
 2.76%    103.01ms        8    12.876ms    8.8060us    102.92ms    cudaStreamCreate
 0.22%    8.2736ms        4    2.0684ms    2.0318ms    2.1029ms    cudaHostAlloc
 0.03%    1.0010ms       182    5.4990us    313ns    283.66us    cuDeviceGetAttribute
 0.01%    418.73us        3    139.58us    112.93us    182.84us    cudaMalloc
 0.01%    196.36us        8    24.544us    19.406us    37.244us    cudaLaunch
 0.00%    150.55us        2    75.273us    70.364us    80.183us    cuDeviceTotalMem
 0.00%    71.331us        2    35.665us    35.460us    35.871us    cuDeviceGetName
 0.00%    58.397us        8    7.2990us    3.7890us    16.599us    cudaMemcpyAsync
 0.00%    9.7850us       32    305ns    118ns    3.4080us    cudaSetupArgument
 0.00%    8.0490us        8    1.0060us    892ns    1.1390us    cudaConfigureCall
 0.00%    2.6620us        1    2.6620us    2.6620us    2.6620us    cudaHostGetDevicePointer
 0.00%    2.4260us        6    404ns    331ns    513ns    cuDeviceGet
 0.00%    2.1840us        3    728ns    308ns    1.4760us    cuDeviceGetCount

```

Figure 3.2. Summary of execution time for each activity on GPU processing matrix multiplication with breadth-approach